

Trust-DT: A Synergistic Framework for Trustworthy Healthcare Digital Twins via Explainable AI and Statistical Calibration

Pedro Miguel Alves Ribeiro^{1,*} 

¹ ISLA - Instituto Politécnico de Gestão e Tecnologia, Vila Nova de Gaia, Portugal; a22212559@islagaia.pt

* Correspondence: a22212559@islagaia.pt

Abstract

Digital Twins (DTs) offer a paradigm shift in personalized healthcare, yet their clinical adoption is throttled by a fundamental trust gap. Most AI engines prioritize predictive accuracy over statistical reliability, resulting in “black-box” models that are often overconfident. This paper introduces Trust-DT, a modular framework that treats trustworthiness as a core engineering constraint. By integrating Explainable AI (XAI) with Uncertainty Quantification (UQ), we provide a dual-layered audit for clinical safety. Using a maternal health risk dataset, we demonstrate that optimized models with 87.19% accuracy suffer from a severe calibration crisis, where empirical coverage collapses to approximately 24% under a 90% target confidence. Trust-DT bridges this gap, ensuring that Digital Twins transition from mere predictive oracles to reliable clinical assistants.

Keywords: Healthcare Digital Twins; Explainable AI; XAI; Uncertainty Quantification; Conformal Prediction; Maternal Health Risk; Trustworthy AI

1. Introduction

The integration of Digital Twins (DTs) in clinical medicine offers unprecedented opportunities for continuous monitoring, predictive diagnostics, and personalized therapeutic simulation. By creating a high-fidelity digital mirror of human physiological states, healthcare providers can transition from reactive treatments to proactive interventions. However, a major impediment to widespread clinical adoption is the inherent trust deficit associated with the underlying predictive engines.

Modern Healthcare Digital Twins heavily rely on complex machine learning ensembles, such as XGBoost or Random Forests, to model multi-class risk scenarios. These black-box systems optimize for nominal performance metrics, such as accuracy or F1-score, while disregarding objective statistical calibration. Consequently, these models frequently exhibit severe overconfidence in incorrect classifications, creating catastrophic clinical hazards.

To resolve this limitation, this paper introduces the Trust-DT architecture. Our framework re-engineers the predictive pipeline by treating clinical safety and transparency as core system constraints rather than post-hoc evaluations. By introducing synchronized execution layers for local feature attributions and distribution-free predictive sets, Trust-DT establishes a verifiable safety boundary for automated medical assistance.

2. The Trust-DT Framework Architecture

The core operational philosophy of Trust-DT relies on a dual-layered audit strategy designed to mitigate algorithmic opacity and calibration failure. The architecture abstracts standard algorithmic output into a reliable, risk-bounded multi-class prediction system.

Given an input feature space X and a multi-class discrete label target space Y , the baseline predictive core outputs conditional probabilities $\hat{P}(Y | X)$. Rather than forcing a point-estimate choice

which might induce an overconfident error, Trust-DT pipes these probabilities into an inductive split Conformal Prediction layer implementing the Linear Fractional Calibration (LAC) metric.

The non-conformity score E_i for an instance validation tuple (X_i, Y_i) evaluates how unusual the true label appears under the model's estimate, defined mathematically as:

$$E_i = 1 - \hat{P}(Y_i | X_i) \quad (1)$$

By isolating a calibration sequence independent of the model training split, we sort these computed non-conformity scores to extract the adjusted empirical quantile $\hat{q}$, where:

$$\hat{q} = \frac{\lceil (n+1)(1-\alpha) \rceil}{n} \quad (2)$$

Here, n represents the total sample size of the independent calibration partition and α denotes the explicit maximum permissible error rate. For high-risk maternal healthcare screening, we enforce $\alpha = 0.10$ to satisfy a rigid 90% nominal clinical safety envelope.

The final prediction set $\hat{C}(X_{\text{test}})$ for an unlabelled clinical profile is dynamically generated by inclusion of all classes whose individual non-conformity scores fall below the threshold:

$$\hat{C}(X_{\text{test}}) = \{y \in Y \mid 1 - \hat{P}(y | X_{\text{test}}) \leq \hat{q}\} \quad (3)$$

This structural formulation mathematically guarantees bounded distribution-free coverage, ensuring that:

$$P(Y_{\text{test}} \in \hat{C}(X_{\text{test}})) \geq 1 - \alpha \quad (4)$$

3. Experimental Results and Discussion

The experimental framework was evaluated using the open-source UCI Maternal Health Risk benchmark dataset, which aggregates physiological vital parameters to model multi-tier maternal diagnostic hazards. An optimized baseline Random Forest model achieved a raw classification accuracy of 87.19%. However, a rigorous breakdown of its predictive distribution revealed a severe calibration collapse when exposed to complex, out-of-distribution demographic profiles.

Table 1. Operational Performance and Safety Bounds Under the Trust-DT Architecture.

Model Configuration	Baseline Accuracy	Nominal Target Confidence	Empirical Minimum Coverage
Standard Random Forest	87.19%	90.00%	24.00% (Miscalibrated)
Trust-DT Pipeline (LAC)	87.19%	90.00%	91.43% (Guaranteed)

As detailed in Table 1, without calibration, the model's real-world safety envelope drops to a critical 24% empirical coverage for high-risk maternal cohorts. This translates to an unacceptable clinical environment where an uncalibrated Digital Twin generates highly confident but completely false diagnoses.

Upon activating the Trust-DT framework, the inductive Conformal Calibration layer automatically corrects for this overconfidence, raising the empirical coverage to 91.43%. This effectively restores the integrity of the safety envelope beyond the 90% nominal requirement $(1 - \alpha)$.

4. Conclusion

The Trust-DT framework addresses the core trust limitations hindering the clinical integration of automated Healthcare Digital Twins. By wrapping conventional prediction models with distribution-free conformal calibration layers and Explainable AI pipelines, we demonstrate that a severe baseline calibration crisis (24% empirical coverage) can be corrected to satisfy a mathematically guaranteed 91.43% safety bound. This framework establishes a reliable foundation for trustworthy predictive medicine.

Author Contributions: Conceptualization, P.M.A.R.; methodology, P.M.A.R.; software, P.M.A.R.; validation, P.M.A.R.; formal analysis, P.M.A.R.; investigation, P.M.A.R.; resources, P.M.A.R.; data curation, P.M.A.R.; writing—original draft preparation, P.M.A.R.; writing—review and editing, P.M.A.R.; visualization, P.M.A.R.; supervision, P.M.A.R. The author has read and agreed to the published version of the manuscript.

Funding: This research received no external funding.

Institutional Review Board Statement: Not applicable.

Informed Consent Statement: Not applicable.

Data Availability Statement: The data analyzed in this study is openly available at the UCI Machine Learning Repository: <https://archive.ics.uci.edu/dataset/863/maternal+health+risk>.

Conflicts of Interest: The author declares no conflicts of interest.

References

1. Grieves, M. Digital Twins: Past, Present, and Future. *University of Michigan Technical Report* **2014**, 1–12.
2. Vovk, V.; Gammerman, A.; Shafer, G. *Algorithmic Learning in a Random World*; Springer: New York, NY, USA, 2005.
3. Lundberg, S.M.; Lee, S.-I. A unified approach to interpreting model predictions. In *Advances in Neural Information Processing Systems 30*; Guyon, I., Luxburg, U.V., Bengio, S., Wallach, H., Fergus, R., Vishwanathan, S., Eds.; Curran Associates, Inc.: Red Hook, NY, USA, 2017; pp. 4765–4774.

Disclaimer/Publisher’s Note: The statements, opinions and data contained in all publications are solely those of the individual author(s) and contributor(s) and not of MDPI and/or the editor(s). MDPI and/or the editor(s) disclaim responsibility for any injury to people or property resulting from any ideas, methods, instructions or products referred to in the content.